

10601a: Homework #5 - “Decision Trees”

TA: Joe Runde (jrunde@andrew.cmu.edu)

Assigned: Monday Feb 8, 2016

Due: 11:59pm on Monday Feb 15

Late Penalty: 20% per day.

Office Hours:

	GHC Citadel Commons
Tues 2/9	4pm
Wed 2/10	4pm
Thurs 2/11	5pm
Sun 2/14	2pm
Mon 2/15	5pm

0 POLICY ON COLLABORATION AMONG STUDENTS

0.1 PREVIOUSLY USED ASSIGNMENTS

Some of the homework assignments used in this class may have been used in prior versions of this class, or in classes at other institutions, or elsewhere. Avoiding the use of heavily tested assignments will detract from the main purpose of these assignments, which is to reinforce the material and stimulate thinking. Because some of these assignments may have been used before, solutions to them may be, or may have been, available online, or from other people or sources. It is explicitly forbidden to use any such sources, or to consult people who have solved these problems before. It is explicitly forbidden to search for these problems or their solutions on the internet. You must solve the homework assignments completely on your own. **For programming assignments, this means you must write your programs completely by yourself, and not use any code from any source whatsoever.** I will be actively monitoring your compliance, and any violation will be dealt with harshly. Collaboration with other students who are currently taking the class is allowed, but only under the conditions stated below.

0.2 COLLABORATION AMONG STUDENTS

The purpose of student collaboration is to facilitate learning, not to circumvent it. Studying the material in groups is strongly encouraged. It is also allowed to seek help from other students in understanding the material needed to solve a particular homework problem, provided no written notes are shared, or are taken at that time, and provided learning is facilitated, not circumvented. **The actual solution must be done by each student alone.**

The purpose of programming assignments in this course is to make sure you truly understand the relevant techniques. In my 20 years of teaching, I have found no better way to achieve this than by having each student struggle by him/herself to implement these techniques “from scratch”. For this reason, in the case of programming assignments **all code must be written by each student alone.** We will strictly enforce this policy, by carefully inspecting your code using sophisticated detection techniques. You have been warned!

The presence or absence of any form of help or collaboration, whether given or received, must be explicitly stated and disclosed in full by all involved. Specifically, each assignment solution must include answering the following questions:

- Did you receive any help whatsoever from anyone in solving this assignment? Yes / No. If you answered 'yes', give full details: (e.g. "Jane Doe explained to me what is asked in Question 3.4")
- Did you give any help whatsoever to anyone in solving this assignment? Yes / No. If you answered 'yes', give full details: (e.g. "I pointed Joe Smith to section 2.3 since he didn't know how to proceed with Question 2")
- Did you find or come across code that implements any part of this assignment ? Yes / No. (See below policy on "found code") If you answered 'yes', give full details: (book & page, URL & location within the page, etc.).

If you gave help after turning in your own assignment and/or after answering the questions above, you must update your answers before the assignment's deadline, if necessary by emailing the TA in

charge of the assignment. Collaboration without full disclosure will be handled severely, in compliance with CMU's Policy on Cheating and Plagiarism.

0.3 POLICY REGARDING "FOUND CODE":

You are encouraged to read books and other instructional materials, both online and offline, to help you understand the concepts and algorithms taught in class. These materials may contain example code or pseudo code, which may help you better understand an algorithm or an implementation detail. However, when you implement your own solution to an assignment, you must put all materials aside, and write your code **completely on your own, starting "from scratch"**. Specifically, you may not use any code you found or came across. If you find or come across code that implements any part of your assignment, you must disclose this fact in your collaboration statement.

0.4 DUTY TO PROTECT ONE'S WORK

Students are responsible for pro-actively protecting their work from copying and misuse by other students. If a student's work is copied by another student, the original author is also considered to be at fault and in gross violation of the course policies. It does not matter whether the author allowed the work to be copied or was merely negligent in preventing it from being copied. When overlapping work is submitted by different students, both students will be punished.

To protect future students, do not post your solutions publicly, neither during the course nor afterwards.

0.5 SEVERE PUNISHMENT OF VIOLATIONS OF COURSE POLICIES

All violations (even first one) of course policies will **always** be reported to the university authorities, will carry **severe** penalties, usually **failure** in the course, and can even lead to **dismissal** from the university. This is not an idle threat - it is my standard practice. You have been warned!

1 GENERAL INSTRUCTIONS

The goal of this assignment is for you to implement a **decision tree learner**, entirely from scratch. We're going to try decision trees in two domains. For simplicity, all variables are discretized into just **two categories**. The datasets for this assignment are available on autolab, go to Homework 5 and click on "Download Handout."

The first task is to predict whether a US **politician** is a member of the Democrat or Republican party, based on their past voting history. Attributes (covariates, predictors) are short descriptions of bills that were voted on, such as **Aid_to_nicaraguan_contras** or **Duty_free_exports**. Values are given as 'y' for yes votes and 'n' for no votes. Check the .csv files to see more.

The second task is to predict the final **grade** (A, not A) for high school students. The attributes (covariates, predictors) are student grades on 5 multiple choice assignments **M1** through **M5**, 4 programming assignments **P1** through **P4**, and the final exam **F**. Check the csv files to see the attribute values.

Before you begin, think about whether decision trees are appropriate for these two tasks. Write down your thoughts in a text file **Q0.txt**: do you think a decision tree will work for the politician dataset? What about for the education dataset?

Next, inspect the training data manually; look for any unusual findings and think about which variables seem useful. In **Q1_politicians.txt** make your best guess about which variables are useful for the politician task. Do the same in **Q1_education.txt**.

We've provided you with attributes and labels split into training and testing data in files "politicians*.csv" and "education*.csv". Throughout, we show results for "example1.csv" and "example2.csv," a small, purely for demonstration version of the politicians dataset, with attributes **Anti_satellite_test_ban** and **Export_south_africa**. The format is comma separated, one row per observation, and one column per attribute. The first line of the file contains the name of each attribute, and the class is always the last column. **Your program will take a training and a testing dataset as input.**

2 WARMUP

First, let's think a little bit about decision trees. Note that, throughout this homework, we will use the convention that the leaves of the trees count as nodes, and as such are included in calculations of depth and number of splits. (For example, a tree which classifies the data based on the value of a single attribute will have depth 1, and contain 1 split.) Please answer the following questions by writing in text files:

X	Y	Z	Class
1	1	1	0
1	1	0	1
1	0	1	1
1	0	0	0
0	1	1	0
0	0	0	0

Q2.txt (2 pts.) Consider the dataset given above. Which attribute (X, Y, or Z) would a decision tree algorithm pick first to branch on, if its criteria is Mutual Information? (Please include only one character in your file)

Q3.txt (1 pts.) If the same algorithm continues until the tree is consistent with the data, what would the depth of the tree be? (Include only one number in your file)

Q4.txt (1 pts.) What depth is the shortest decision tree which is consistent with the data? (Include only one number in your file)

Now on to programming, you may use Java or python to complete this assignment. Make sure to use the versions of python and Java supported by Autolab (currently python 2.6, Java 7), or your submission will not be graded correctly. Write a program `inspect.py` or `inspect.java` to calculate the label `entropy` at the root (i.e. the entropy of the labels before any splits) and the `error` rate (the percent of incorrectly classified instances) of classifying using a `majority` vote (picking the label with the most examples). You do not need to look at the values of any of the attributes to do these calculations, knowing the labels of each example is sufficient.

```
$ python inspect.py example1.csv
entropy: 0.996316519559
error: 0.464285714286
```

```
$ javac inspect.java
$ java inspect example1.csv
entropy: 0.996316519559
error: 0.464285714286
```

Test your program on `both datasets`—this error rate is a baseline over which we would (ideally) like to improve.

3 TRAINING THE TREE

Implement a decision tree learner with the following guidelines. (As a reference, consult Mitchell, Chapter 3, page 56.)

- Use **mutual information** to determine which attribute to split on
- Be sure you're correctly **weighting** your calculation of mutual information. For a split on attribute X , $I(Y; X) = H(Y) - H(Y|X) = H(Y) - P(X=0)H(Y|X=0) - P(X=1)H(Y|X=1)$. Equivalently, you can calculate $I(Y; X) = H(Y) + H(X) - H(Y, X)$.
- As a **stopping rule**, only split on an attribute if the mutual information is $\geq .1$.
- Do not grow the tree **beyond depth 2**. Namely, split a node only if the mutual information is $\geq .1$ and the node is the **root or a direct child of the root**.
- Use a **majority vote** of the labels at each leaf to make classification decisions

Hints on getting started: write helper functions to calculate entropy and mutual information. Write a function to train a stump (tree with only one level). The correct tree and output format for the example data are shown below, where we are training on example1.csv and testing on example2.csv. For the politician data, use "+" for Party = "democrat" and "-" for Party = "republican". With the education data, use "+" for final **grade** = "A" and "-" for final **grade** = "not A". Don't worry about the order in which you list the left and right children, the autograder will take care of it. Your program should be named decisionTree and take two arguments, a training file, and a test file.

```
$ python decisionTree.py example1.csv example2.csv
or
$ javac decisionTree.java
$ java decisionTree example1.csv example2.csv
```

```
[15+/13-]
Anti_satellite_test_ban = y: [13+/1-]
| Export_south_africa = y: [13+/0-]
| Export_south_africa = n: [0+/1-]
Anti_satellite_test_ban = n: [2+/12-]
| Export_south_africa = y: [2+/7-]
| Export_south_africa = n: [0+/5-]
error(train): 0.0714285714286
error(test): 0.142857142857
```

Remember, the tree might not be full. Here's what happens when we train on example2.csv and test on example1.csv:

```
$ python decisionTree.py example2.csv example1.csv
or
$ javac decisionTree.java
$ java decisionTree example2.csv example1.csv
```

```
[13+/15-]
Anti_satellite_test_ban = y: [9+/0-]
Anti_satellite_test_ban = n: [4+/15-]
| Export_south_africa = y: [4+/10-]
| Export_south_africa = n: [0+/5-]
error(train): 0.142857142857
error(test): 0.107142857143
```

The numbers in brackets give the number of positive and negative labels from the training data in that part of the tree. The last two numbers are the error rate on the training data and the error rate on the testing data. Make sure your answers are accurate to within 0.01.

4 EVALUATION

Train and test a decision tree for the politician dataset and the education dataset. Which is more accurate on the training data? Which is more accurate on the testing data? Write down your observations in **Q5.txt**.

In addition to the politician and education datasets, autolab will test your code on two more datasets, which will not be shown to you. One set contains information about various cars, and whether or not consumers decided to buy them. The other contains data about songs, and whether or not they became top hits. The data will be in .csv files similar to the ones provided, again with the class as the last column. Shown below are samples of the attributes and the values they can take:

Music data:

- Attribute: year ('before1950' or 'after1950')
- Attribute: solo ('yes' or 'no')
- Attribute: vocal ('yes' or 'no')
- Attribute: length ('morethan3min' or 'lessthan3min')
- Attribute: original ('yes' or 'no')
- Attribute: tempo ('fast' or 'slow')
- Attribute: folk ('yes' or 'no')
- Attribute: classical ('yes' or 'no')
- Attribute: rhythm ('yes' or 'no')
- Attribute: jazz ('yes' or 'no')
- Attribute: rock ('yes' or 'no')
- Class Label: hit ('yes' or 'no')

Cars data:

- Attribute: buying ('expensive' or 'cheap')
- Attribute: maint ('high' or 'low')
- Attribute: doors ('Two' or 'MoreThanTwo')
- Attribute: person ('Two' or 'MoreThanTwo')
- Attribute: boot ('large' or 'small')
- Attribute: safety ('high' or 'low')
- Class Label: class ('yes' or 'no')

Please ensure your solution can handle data with these values, and use "yes" for "+".

5 AUTOLAB SUBMISSION

Submit a .tgz file containing your source code, written assignments, and collaboration.txt. You can create this archive by running `tar -cvf hw5.tgz *.py *.java *.txt`. **DO NOT** put the above files in a folder and then tar gzip the folder. You must submit this file to the "homework5" link on Autolab. There will be no limit on the number of submissions.